

```
In [1]: import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader
import torchvision
import torchvision.transforms as transforms
```

Creating Tensors

```
In [2]: torch.empty(2), torch.ones(2), torch.zeros(2)
```

```
Out[2]: (tensor([1.1422e-40, 4.5703e-35]), tensor([1., 1.]), tensor([0., 0.]))
```

```
In [3]: torch.rand(5), torch.randn(5)
```

```
Out[3]: (tensor([0.9261, 0.7819, 0.3198, 0.2448, 0.3161]),
tensor([ 0.0261, -1.8937, -1.3276,  2.1708, -1.8366]))
```

```
In [4]: t = torch.tensor(1)
t
```

```
Out[4]: tensor(1)
```

```
In [5]: t = torch.tensor([1])
t
```

```
Out[5]: tensor([1])
```

```
In [6]: torch.tensor([1, 2, 3]).dtype
```

```
Out[6]: torch.int64
```

```
In [7]: torch.tensor([1.0, 2.0, 3.0]).dtype
```

```
Out[7]: torch.float32
```

```
In [8]: torch.tensor([1], dtype=torch.int32).dtype
```

```
Out[8]: torch.int32
```

```
In [9]: v = np.array([1, 2, 3])
t = torch.from_numpy(v)
t
```

```
Out[9]: tensor([1, 2, 3])
```

```
In [10]: v += 1
t
```

```
Out[10]: tensor([2, 3, 4])
```

```
In [11]: u = t.numpy()
u
```

```
Out[11]: array([2, 3, 4])
```

```
In [12]: t.add_(1)
         u
```

```
Out[12]: array([3, 4, 5])
```

```
In [13]: t = torch.tensor(range(1,7)).view(-1, 3)
         t
```

```
Out[13]: tensor([[1, 2, 3],
                 [4, 5, 6]])
```

Indexing & Slicing

```
In [14]: t[0]
```

```
Out[14]: tensor([1, 2, 3])
```

```
In [15]: t[0, 1]
```

```
Out[15]: tensor(2)
```

```
In [16]: s = t[0, 1].item()
         s, type(s)
```

```
Out[16]: (2, int)
```

```
In [17]: t[:, 2], t[1, :]
```

```
Out[17]: (tensor([3, 6]), tensor([4, 5, 6]))
```

Arithmetic

```
In [18]: t + t
```

```
Out[18]: tensor([[ 2,  4,  6],
                 [ 8, 10, 12]])
```

```
In [19]: t * t
```

```
Out[19]: tensor([[ 1,  4,  9],
                 [16, 25, 36]])
```

```
In [20]: t ** 2
```

```
Out[20]: tensor([[ 1,  4,  9],
                 [16, 25, 36]])
```

```
In [21]: t @ t.T # 2x3 . 3x2 = 2x2
```

```
Out[21]: tensor([[14, 32],
                 [32, 77]])
```

```
In [22]: t.add(2)
```

```
Out[22]: tensor([[3, 4, 5],
                [6, 7, 8]])
```

```
In [23]: t.add_(2), t
```

```
Out[23]: (tensor([[3, 4, 5],
                [6, 7, 8]]),
          tensor([[3, 4, 5],
                [6, 7, 8]]))
```

Aggregation

```
In [24]: t
```

```
Out[24]: tensor([[3, 4, 5],
                [6, 7, 8]])
```

```
In [25]: t.max(), torch.max(t)
```

```
Out[25]: (tensor(8), tensor(8))
```

```
In [26]: t.max(dim=0), t.min(dim=0)
```

```
Out[26]: (torch.return_types.max(
  values=tensor([6, 7, 8]),
  indices=tensor([1, 1, 1])),
  torch.return_types.min(
  values=tensor([3, 4, 5]),
  indices=tensor([0, 0, 0])))
```

```
In [27]: torch.max(t, dim=1).values, torch.max(dim=1).indices
```

```
Out[27]: (tensor([5, 8]), tensor([2, 2]))
```

```
In [28]: t.argmax(dim=1), t.argmax(dim=1, keepdim=True)
```

```
Out[28]: (tensor([2, 2]),
          tensor([[2],
                [2]]))
```

AutoGrad

```
In [29]: t = torch.tensor(1.0, requires_grad=True)
          t
```

```
Out[29]: tensor(1., requires_grad=True)
```

```
In [30]: t.detach(), t
```

```
Out[30]: (tensor(1.), tensor(1., requires_grad=True))
```

```
In [31]: t.requires_grad_(False), t
```

```
Out[31]: (tensor(1.), tensor(1.))
```

```
In [32]: x = torch.tensor(1.5, requires_grad=True)
        y = 3*x**2
        y
```

```
Out[32]: tensor(6.7500, grad_fn=<MulBackward0>)
```

```
In [33]: y.backward()
        x.grad
```

```
Out[33]: tensor(9.)
```

```
In [34]: x.grad.zero_(), x.grad
```

```
Out[34]: (tensor(0.), tensor(0.))
```

```
In [35]: with torch.no_grad():
        z = 0.5*x
        try:
            z.backward()
        except:
            print('Exception')
        z, x.grad
```

```
Exception
Out[35]: (tensor(0.7500), tensor(0.))
```

nn.CrossEntropyLoss(), nn.NLLLoss()

```
In [36]: logits = torch.tensor([[1.0, 2.0, 3.0]])
        probs = torch.softmax(logits, dim=-1)
        probs
```

```
Out[36]: tensor([[0.0900, 0.2447, 0.6652]])
```

```
In [37]: ce_loss = nn.CrossEntropyLoss()
        y0 = torch.tensor([0])
        y2 = torch.tensor([2])
        # CrossEntropyLoss() takes in *logits (NOT probs)* and *class labels*
        ce_loss(logits, y0), ce_loss(logits, y2)
```

```
Out[37]: (tensor(2.4076), tensor(0.4076))
```

```
In [38]: -np.log(probs.numpy())
```

```
Out[38]: array([[2.408, 1.408, 0.408]], dtype=float32)
```

```
In [39]: nll_loss = nn.NLLLoss()
        nll_loss(probs, y0), nll_loss(probs, y2)
```

```
Out[39]: (tensor(-0.0900), tensor(-0.6652))
```

nn.BCELoss()

```
In [40]: bce_loss = nn.BCELoss()
bce_loss(torch.tensor([0.9]), torch.tensor([0.0]))
```

```
Out[40]: tensor(2.3026)
```

```
In [41]: -(1-0)*np.log(1 - 0.9)
```

```
Out[41]: 2.303
```

Backpropagation

```
In [42]: X = torch.tensor([1, 2, 3, 4], dtype=torch.float32).view(-1, 1)
Y = torch.tensor([2, 4, 6, 8], dtype=torch.float32).view(-1, 1)

w = torch.tensor([[0.0]], dtype=torch.float32, requires_grad=True)

def forward(X, w):
    return X @ w # 4x1 . 1x1 = 4x1

def loss(Y_pred, Y):
    return ((Y_pred - Y)**2).mean()

learning_rate = 0.01

for epoch in range(40):
    Y_pred = forward(X, w)
    j = loss(Y, Y_pred)
    j.backward()
    with torch.no_grad():
        w -= learning_rate * w.grad
    w.grad.zero_()

    if epoch % 10 == 9:
        print(f'epoch {epoch+1}: w = {w.item():.3f}, loss = {j.item():.8f}')

epoch 10: w = 1.606, loss = 1.60939169
epoch 20: w = 1.922, loss = 0.06237914
epoch 30: w = 1.985, loss = 0.00241778
epoch 40: w = 1.997, loss = 0.00009371
```

Optimizer & LR-Scheduler

```
In [43]: w = torch.tensor([[0.0]], dtype=torch.float32, requires_grad=True)
optimizer = torch.optim.SGD([w], lr=learning_rate)

loss = nn.MSELoss()

for epoch in range(40):
    Y_pred = forward(X, w)
    j = loss(Y_pred, Y)
    j.backward()
    optimizer.step()
    optimizer.zero_grad()

    if epoch % 10 == 9:
        print(f'epoch {epoch+1}: w = {w.item():.3f}, loss = {j.item():.8f}')
```

```
epoch 10: w = 1.606, loss = 1.60939169
epoch 20: w = 1.922, loss = 0.06237914
epoch 30: w = 1.985, loss = 0.00241778
epoch 40: w = 1.997, loss = 0.00009371
```

```
In [44]: lr_scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=10, gamma=0.1)

for epoch in range(30):
    Y_pred = forward(X, w)
    j = loss(Y_pred, Y)
    j.backward()
    optimizer.step()
    optimizer.zero_grad()
    lr_scheduler.step()

    if epoch % 10 == 9:
        lr = optimizer.state_dict()['param_groups'][0]['lr']
        print(f'epoch {epoch+1}: w = {w.item():.3f}, lr = {j.item():.8f}')

epoch 10: w = 1.999, lr = 0.00000363
epoch 20: w = 2.000, lr = 0.00000064
epoch 30: w = 2.000, lr = 0.00000028
```

nn.Linear()

```
In [45]: model = nn.Linear(in_features=1, out_features=1)
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

for epoch in range(40):
    Y_pred = model(X)
    j = loss(Y_pred, Y)
    j.backward()
    optimizer.step()
    optimizer.zero_grad()

    if epoch % 10 == 9:
        w, b = model.parameters() # unpack parameters
        print(f'epoch {epoch+1}: w = {w.item():.3f}, loss = {j.item():.8f}')

epoch 10: w = 1.948, loss = 0.65129113
epoch 20: w = 2.122, loss = 0.05157164
epoch 30: w = 2.147, loss = 0.03403451
epoch 40: w = 2.147, loss = 0.03167749
```

```
In [46]: predicted = model(X).detach()
predicted
```

```
Out[46]: tensor([[1.7123],
                [3.8591],
                [6.0059],
                [8.1526]])
```

```
In [47]: with torch.no_grad():
    predicted = model(X)
predicted
```

```
Out[47]: tensor([[1.7123],
                [3.8591],
                [6.0059],
                [8.1526]])
```

nn.Module, Logistic Regression

```
In [48]: X = torch.tensor([[2, 1], [1, 2]], dtype=torch.float32)
Y = torch.tensor([0, 1], dtype=torch.float32).view(-1, 1)
```

```
class LogR(nn.Module):
    def __init__(self, in_features):
        super().__init__()
        self.linear = nn.Linear(in_features, 1)

    def forward(self, x):
        x = self.linear(x)
        x = torch.sigmoid(x)
        return x

model = LogR(X.shape[-1])
optimizer = torch.optim.SGD(model.parameters(), lr=0.1)
loss = nn.BCELoss()

for epoch in range(50):
    Y_pred = model(X)
    j = loss(Y_pred, Y)
    j.backward()
    optimizer.step()
    optimizer.zero_grad()

    if epoch % 10 == 9:
        print(f'epoch {epoch+1}: loss = {j.item():.8f}')
```

```
epoch 10: loss = 0.58357650
epoch 20: loss = 0.47398025
epoch 30: loss = 0.40872616
epoch 40: loss = 0.35784590
epoch 50: loss = 0.31676430
```

```
In [49]: with torch.no_grad():
        Y_pred = model(X)
        Y_pred
```

```
Out[49]: tensor([[0.2693],
                [0.7317]])
```

Misc

```
In [50]: device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
torch.tensor([1], device=device), torch.tensor([2]).to(device)
```

```
Out[50]: (tensor([1]), tensor([2]))
```

Model Saving & Loading

```
In [51]: class Model(nn.Module):
        def __init__(self):
            super().__init__()
```

```

        self.linear = nn.Linear(2, 1)

    def forward(self, x):
        x = self.linear(x)
        x = torch.sigmoid(x)
        return x

```

```

In [52]: model = Model()
torch.save(model, 'model.pth')
# ---
model = torch.load('model.pth')
model.eval()
for param in model.parameters():
    print(param)

```

```

Parameter containing:
tensor([[ -0.0505,  0.6279]], requires_grad=True)
Parameter containing:
tensor([0.0782], requires_grad=True)

```

```

In [53]: model = Model()
torch.save(model.state_dict(), 'model_state.pth')
print(model.state_dict())
# ---
model = Model()
model.load_state_dict(torch.load('model_state.pth'))
# model.load_state_dict(torch.load('model_state.path', map_location=device))
model.eval()
print(model.state_dict())

```

```

OrderedDict([('linear.weight', tensor([[ 0.3975, -0.2836]])), ('linear.bias',
tensor([0.2074]))])
OrderedDict([('linear.weight', tensor([[ 0.3975, -0.2836]])), ('linear.bias',
tensor([0.2074]))])

```

```

In [54]: checkpoint = {
    "model_state": model.state_dict(),
    "optim_state": optimizer.state_dict()
}
torch.save(checkpoint, 'checkpoint.pth')
# ---
model = Model()
optimizer = torch.optim.SGD(model.parameters(), lr=0.01)
checkpoint = torch.load('checkpoint.pth')
model.load_state_dict(checkpoint['model_state'])
optimizer.load_state_dict(checkpoint['optim_state'])
optimizer

```

```

Out[54]: SGD (
  Parameter Group 0
    dampening: 0
    lr: 0.1
    momentum: 0
    nesterov: False
    weight_decay: 0
)

```

Dataset & DataLoader


```
In [55]: class MyDataset(Dataset):
        def __init__(self):
            super().__init__()
            self.X = torch.tensor([[2, 1], [1, 2]], dtype=torch.float32)
            self.Y = torch.tensor([0, 1], dtype=torch.float32).view(-1, 1)

        def __getitem__(self, idx):
            return self.X[idx], self.Y[idx]

        def __len__(self):
            return len(self.X)

ds = MyDataset()
```

```
In [56]: len(ds), ds[0]
```

```
Out[56]: (2, (tensor([2., 1.]), tensor([0.])))
```

```
In [57]: dl = DataLoader(ds, batch_size=2, shuffle=True)
        next(iter(dl))
```

```
Out[57]: [tensor([[2., 1.],
                [1., 2.])),
          tensor([[0.],
                [1.]])]
```

```
In [58]: for epoch in range(50):
        for X, Y in dl:
            X = X.to(device)
            Y = Y.to(device)
            Y_pred = model(X)
            j = loss(Y_pred, Y)
            optimizer.zero_grad()
            j.backward()
            optimizer.step()

        if epoch % 10 == 9:
            print(f'epoch {epoch+1}: loss = {j.item():.8f}')
```

```
epoch 10: loss = 0.74077058
epoch 20: loss = 0.61800981
epoch 30: loss = 0.52327919
epoch 40: loss = 0.44903362
epoch 50: loss = 0.39019936
```

```
In [59]: # ds = torchvision.datasets.ImageFolder(path)
```

Torchvision Transforms

```
In [60]: transform = transforms.Compose([transforms.RandomHorizontalFlip(),
                                       transforms.ToTensor(),
                                       transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
```

```
In [61]: train_dataset = torchvision.datasets.CIFAR10(
        root='~/pytorch_datasets', download=True, train=True, transform=transform)

Files already downloaded and verified
```

```
In [62]: train_dataset[0][0].shape # NOTE: PyTorch uses Channel-FIRST (TF uses Channel-
```

```
Out[62]: torch.Size([3, 32, 32])
```

Transfer Learning

```
In [63]: model = torchvision.models.resnet18(pretrained=True)
model.fc = nn.Linear(model.fc.in_features, 10)
for param in model.parameters():
    param.requires_grad = False
```

```
In [ ]:
```