

# PyTorch Cheat Sheet

```
In [1]: import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader
import torchvision
import torchvision.transforms as transforms
```

## Creating Tensors

```
In [2]: torch.empty(2), torch.ones(2), torch.zeros(2)
Out[2]: (tensor([1.1422e-40, 4.5703e-35]), tensor([1., 1.]), tensor([0., 0.]))
```

```
In [3]: torch.rand(5), torch.randn(5)
Out[3]: (tensor([0.9261, 0.7819, 0.3198, 0.2448, 0.3161]),
tensor([ 0.0261, -1.8937, -1.3276,  2.1708, -1.8366]))
```

```
In [4]: t = torch.tensor(1)
t
```

```
Out[4]: tensor(1)
```

```
In [5]: t = torch.tensor([1])
t
```

```
Out[5]: tensor([1])
```

```
In [6]: torch.tensor([1, 2, 3]).dtype
```

```
Out[6]: torch.int64
```

```
In [7]: torch.tensor([1.0, 2.0, 3.0]).dtype
```

```
Out[7]: torch.float32
```

```
In [8]: torch.tensor([1], dtype=torch.int32).dtype
```

```
Out[8]: torch.int32
```

```
In [9]: v = np.array([1, 2, 3])
t = torch.from_numpy(v)
t
```

```
Out[9]: tensor([1, 2, 3])
```

```
In [10]: v += 1
t
```

```
Out[10]: tensor([2, 3, 4])
```

```
In [11]: u = t.numpy()  
u
```

```
Out[11]: array([2, 3, 4])
```

```
In [12]: t.add_(1)  
u
```

```
Out[12]: array([3, 4, 5])
```

```
In [13]: t = torch.tensor(range(1,7)).view(-1, 3)  
t
```

```
Out[13]: tensor([[1, 2, 3],  
                [4, 5, 6]])
```

## Indexing & Slicing

```
In [14]: t[0]
```

```
Out[14]: tensor([1, 2, 3])
```

```
In [15]: t[0, 1]
```

```
Out[15]: tensor(2)
```

```
In [16]: s = t[0, 1].item()  
s, type(s)
```

```
Out[16]: (2, int)
```

```
In [17]: t[:, 2], t[1, :]
```

```
Out[17]: (tensor([3, 6]), tensor([4, 5, 6]))
```

## Arithmetic

```
In [18]: t + t
```

```
Out[18]: tensor([[ 2,  4,  6],  
                [ 8, 10, 12]])
```

```
In [19]: t * t
```

```
Out[19]: tensor([[ 1,  4,  9],  
                [16, 25, 36]])
```

```
In [20]: t ** 2
```

```
Out[20]: tensor([[ 1,  4,  9],  
                [16, 25, 36]])
```

```
In [21]: t @ t.T # 2x3 . 3x2 = 2x2
```

```
Out[21]: tensor([[14, 32],  
               [32, 77]])
```

```
In [22]: t.add(2)
```

```
Out[22]: tensor([[3, 4, 5],  
               [6, 7, 8]])
```

```
In [23]: t.add_(2), t
```

```
Out[23]: (tensor([[3, 4, 5],  
               [6, 7, 8]]),  
          tensor([[3, 4, 5],  
               [6, 7, 8]]))
```

## Aggregation

```
In [24]: t
```

```
Out[24]: tensor([[3, 4, 5],  
               [6, 7, 8]])
```

```
In [25]: t.max(), torch.max(t)
```

```
Out[25]: (tensor(8), tensor(8))
```

```
In [26]: t.max(dim=0), t.min(dim=0)
```

```
Out[26]: (torch.return_types.max(  
          values=tensor([6, 7, 8]),  
          indices=tensor([1, 1, 1])),  
          torch.return_types.min(  
          values=tensor([3, 4, 5]),  
          indices=tensor([0, 0, 0])))
```

```
In [27]: torch.max(t, dim=1).values, torch.max(dim=1).indices
```

```
Out[27]: (tensor([5, 8]), tensor([2, 2]))
```

```
In [28]: t.argmax(dim=1), torch.argmax(dim=1, keepdim=True)
```

```
Out[28]: (tensor([2, 2]),  
          tensor([[2],  
                [2]]))
```

## AutoGrad

```
In [29]: t = torch.tensor(1.0, requires_grad=True)  
t
```

```
Out[29]: tensor(1., requires_grad=True)
```

```
In [30]: t.detach(), t
```

```
Out[30]: (tensor(1.), tensor(1., requires_grad=True))
```

```
In [31]: t.requires_grad_(False), t
```

```
Out[31]: (tensor(1.), tensor(1.))
```

```
In [32]: x = torch.tensor(1.5, requires_grad=True)
y = 3*x**2
y
```

```
Out[32]: tensor(6.7500, grad_fn=<MulBackward0>)
```

```
In [33]: y.backward()
x.grad
```

```
Out[33]: tensor(9.)
```

```
In [34]: x.grad.zero_(), x.grad
```

```
Out[34]: (tensor(0.), tensor(0.))
```

```
In [35]: with torch.no_grad():
          z = 0.5*x
        try:
            z.backward()
        except:
            print('Exception')
z, x.grad
```

```
Exception
Out[35]: (tensor(0.7500), tensor(0.))
```

## nn.CrossEntropyLoss(), nn.NLLLoss()

```
In [36]: logits = torch.tensor([[1.0, 2.0, 3.0]])
probs = torch.softmax(logits, dim=-1)
probs
```

```
Out[36]: tensor([[0.0900, 0.2447, 0.6652]])
```

```
In [37]: ce_loss = nn.CrossEntropyLoss()
y0 = torch.tensor([0])
y2 = torch.tensor([2])
# CrossEntropyLoss() takes in *logits (NOT probs)* and *class labels*
ce_loss(logits, y0), ce_loss(logits, y2)
```

```
Out[37]: (tensor(2.4076), tensor(0.4076))
```

```
In [38]: -np.log(probs.numpy())
```

```
Out[38]: array([[2.408, 1.408, 0.408]], dtype=float32)
```

```
In [39]: nll_loss = nn.NLLLoss()
nll_loss(probs, y0), nll_loss(probs, y2)
```

```
Out[39]: (tensor(-0.0900), tensor(-0.6652))
```

## nn.BCELoss()

```
In [40]: bce_loss = nn.BCELoss()  
bce_loss(torch.tensor([0.9]), torch.tensor([0.0]))
```

```
Out[40]: tensor(2.3026)
```

```
In [41]: -(1-0)*np.log(1 - 0.9)
```

```
Out[41]: 2.303
```

## Backpropagation

```
In [42]: X = torch.tensor([1, 2, 3, 4], dtype=torch.float32).view(-1, 1)  
Y = torch.tensor([2, 4, 6, 8], dtype=torch.float32).view(-1, 1)  
  
w = torch.tensor([[0.0]], dtype=torch.float32, requires_grad=True)  
  
def forward(X, w):  
    return X @ w # 4x1 . 1x1 = 4x1  
  
def loss(Y_pred, Y):  
    return ((Y_pred - Y)**2).mean()  
  
learning_rate = 0.01  
  
for epoch in range(40):  
    Y_pred = forward(X, w)  
    j = loss(Y, Y_pred)  
    j.backward()  
    with torch.no_grad():  
        w -= learning_rate * w.grad  
    w.grad.zero_()  
  
    if epoch % 10 == 9:  
        print(f'epoch {epoch+1}: w = {w.item():.3f}, loss = {j.item():.8f}')
```

epoch 10: w = 1.606, loss = 1.60939169  
epoch 20: w = 1.922, loss = 0.06237914  
epoch 30: w = 1.985, loss = 0.00241778  
epoch 40: w = 1.997, loss = 0.00009371

## Optimizer & LR-Scheduler

```
In [43]: w = torch.tensor([[0.0]], dtype=torch.float32, requires_grad=True)  
optimizer = torch.optim.SGD([w], lr=learning_rate)  
  
loss = nn.MSELoss()  
  
for epoch in range(40):  
    Y_pred = forward(X, w)
```

```

j = loss(Y_pred, Y)
j.backward()
optimizer.step()
optimizer.zero_grad()

if epoch % 10 == 9:
    print(f'epoch {epoch+1}: w = {w.item():.3f}, loss = {j.item():.8f}')

```

```

epoch 10: w = 1.606, loss = 1.60939169
epoch 20: w = 1.922, loss = 0.06237914
epoch 30: w = 1.985, loss = 0.00241778
epoch 40: w = 1.997, loss = 0.00009371

```

In [44]: `lr_scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=10, gamma=0.1)`

```

for epoch in range(30):
    Y_pred = forward(X, w)
    j = loss(Y_pred, Y)
    j.backward()
    optimizer.step()
    optimizer.zero_grad()
    lr_scheduler.step()

    if epoch % 10 == 9:
        lr = optimizer.state_dict()['param_groups'][0]['lr']
        print(f'epoch {epoch+1}: w = {w.item():.3f}, lr = {j.item():.8f}')

```

```

epoch 10: w = 1.999, lr = 0.00000363
epoch 20: w = 2.000, lr = 0.00000064
epoch 30: w = 2.000, lr = 0.00000028

```

## nn.Linear()

In [45]: `model = nn.Linear(in_features=1, out_features=1)`  
`optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)`

```

for epoch in range(40):
    Y_pred = model(X)
    j = loss(Y_pred, Y)
    j.backward()
    optimizer.step()
    optimizer.zero_grad()

    if epoch % 10 == 9:
        w, b = model.parameters() # unpack parameters
        print(f'epoch {epoch+1}: w = {w.item():.3f}, loss = {j.item():.8f}')

```

```

epoch 10: w = 1.948, loss = 0.65129113
epoch 20: w = 2.122, loss = 0.05157164
epoch 30: w = 2.147, loss = 0.03403451
epoch 40: w = 2.147, loss = 0.03167749

```

In [46]: `predicted = model(X).detach()`  
`predicted`

Out[46]: `tensor([[1.7123],`  
 `[3.8591],`  
 `[6.0059],`  
 `[8.1526]])`

```
In [47]: with torch.no_grad():
         predicted = model(X)
         predicted
```

```
Out[47]: tensor([[1.7123],
                [3.8591],
                [6.0059],
                [8.1526]])
```

## nn.Module, Logistic Regression

```
In [48]: X = torch.tensor([[2, 1], [1, 2]], dtype=torch.float32)
         Y = torch.tensor([0, 1], dtype=torch.float32).view(-1, 1)
```

```
class LogR(nn.Module):
    def __init__(self, in_features):
        super().__init__()
        self.linear = nn.Linear(in_features, 1)

    def forward(self, x):
        x = self.linear(x)
        x = torch.sigmoid(x)
        return x

model = LogR(X.shape[-1])
optimizer = torch.optim.SGD(model.parameters(), lr=0.1)
loss = nn.BCELoss()

for epoch in range(50):
    Y_pred = model(X)
    j = loss(Y_pred, Y)
    j.backward()
    optimizer.step()
    optimizer.zero_grad()

    if epoch % 10 == 9:
        print(f'epoch {epoch+1}: loss = {j.item():.8f}')
```

```
epoch 10: loss = 0.58357650
epoch 20: loss = 0.47398025
epoch 30: loss = 0.40872616
epoch 40: loss = 0.35784590
epoch 50: loss = 0.31676430
```

```
In [49]: with torch.no_grad():
         Y_pred = model(X)
         Y_pred
```

```
Out[49]: tensor([[0.2693],
                [0.7317]])
```

## Misc

```
In [50]: device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
         torch.tensor([1], device=device), torch.tensor([2]).to(device)
```

```
Out[50]: (tensor([1]), tensor([2]))
```

# Model Saving & Loading

```
In [51]: class Model(nn.Module):
          def __init__(self):
              super().__init__()
              self.linear = nn.Linear(2, 1)

          def forward(self, x):
              x = self.linear(x)
              x = torch.sigmoid(x)
              return x
```

```
In [52]: model = Model()
          torch.save(model, 'model.pth')
          # ---
          model = torch.load('model.pth')
          model.eval()
          for param in model.parameters():
              print(param)
```

Parameter containing:  
tensor([[ -0.0505, 0.6279]], requires\_grad=True)  
Parameter containing:  
tensor([0.0782], requires\_grad=True)

```
In [53]: model = Model()
          torch.save(model.state_dict(), 'model_state.pth')
          print(model.state_dict())
          # ---
          model = Model()
          model.load_state_dict(torch.load('model_state.pth'))
          # model.load_state_dict(torch.load('model_state.path', map_location=device))
          model.eval()
          print(model.state_dict())
```

OrderedDict([('linear.weight', tensor([[ 0.3975, -0.2836]])), ('linear.bias', tensor([0.2074]))])  
OrderedDict([('linear.weight', tensor([[ 0.3975, -0.2836]])), ('linear.bias', tensor([0.2074]))])

```
In [54]: checkpoint = {
          "model_state": model.state_dict(),
          "optim_state": optimizer.state_dict()
          }
          torch.save(checkpoint, 'checkpoint.pth')
          # ---
          model = Model()
          optimizer = torch.optim.SGD(model.parameters(), lr=0.01)
          checkpoint = torch.load('checkpoint.pth')
          model.load_state_dict(checkpoint['model_state'])
          optimizer.load_state_dict(checkpoint['optim_state'])
          optimizer
```



```
Out[54]: SGD (
  Parameter Group 0
    dampening: 0
    lr: 0.1
    momentum: 0
    nesterov: False
    weight_decay: 0
)
```

## Dataset & DataLoader

```
In [55]: class MyDataset(Dataset):
    def __init__(self):
        super().__init__()
        self.X = torch.tensor([[2, 1], [1, 2]], dtype=torch.float32)
        self.Y = torch.tensor([0, 1], dtype=torch.float32).view(-1, 1)

    def __getitem__(self, idx):
        return self.X[idx], self.Y[idx]

    def __len__(self):
        return len(self.X)

ds = MyDataset()
```

```
In [56]: len(ds), ds[0]
```

```
Out[56]: (2, (tensor([2., 1.]), tensor([0.])))
```

```
In [57]: dl = DataLoader(ds, batch_size=2, shuffle=True)
next(iter(dl))
```

```
Out[57]: [tensor([[2., 1.],
                [1., 2.])),
          tensor([[0.],
                [1.]])]
```

```
In [58]: for epoch in range(50):
    for X, Y in dl:
        X = X.to(device)
        Y = Y.to(device)
        Y_pred = model(X)
        j = loss(Y_pred, Y)
        optimizer.zero_grad()
        j.backward()
        optimizer.step()

    if epoch % 10 == 9:
        print(f'epoch {epoch+1}: loss = {j.item():.8f}')
```

```
epoch 10: loss = 0.74077058
epoch 20: loss = 0.61800981
epoch 30: loss = 0.52327919
epoch 40: loss = 0.44903362
epoch 50: loss = 0.39019936
```

```
In [59]: # ds = torchvision.datasets.ImageFolder(path)
```

# Torchvision Transforms

```
In [60]: transform = transforms.Compose([transforms.RandomHorizontalFlip(),
                                         transforms.ToTensor(),
                                         transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
```

```
In [61]: train_dataset = torchvision.datasets.CIFAR10(
          root='~/pytorch_datasets', download=True, train=True, transform=transform)
```

Files already downloaded and verified

```
In [62]: train_dataset[0][0].shape # NOTE: PyTorch uses Channel-FIRST (TF uses Channel-LAST)
```

```
Out[62]: torch.Size([3, 32, 32])
```

# Transfer Learning

```
In [63]: model = torchvision.models.resnet18(pretrained=True)
          model.fc = nn.Linear(model.fc.in_features, 10)
          for param in model.parameters():
              param.requires_grad = False
```

```
In [ ]:
```